E09 — Non-overlapping Intervals (Greedy)

Experiment: 9 | **LeetCode:** #435 | **CO:** CO3, CO4 | **BT Level:** BT5

Problem Statement

Given an array of intervals `intervals` where `intervals[i] = [starti, endi]`, return the **minimum number of intervals you need to remove** to make the rest of the intervals non-overlapping.

Note that intervals which only touch at a point are non-overlapping. For example, `[1, 2]` and `[2, 3]` are non-overlapping.

Example 1:

```
Input:  intervals = [[1,2],[2,3],[3,4],[1,3]]
Output: 1
Explanation: Remove [1,3] and the rest are non-overlapping.
```

Example 2:

```
Input:  intervals = [[1,2],[1,2],[1,2]]
Output: 2
Explanation: Remove two [1,2] intervals.
```

Example 3:

```
Input:  intervals = [[1,2],[2,3]]
Output: 0
Explanation: Already non-overlapping (touching at 2 is fine).
```

Concept: Greedy Interval Scheduling

This is equivalent to finding the **maximum number of non-overlapping intervals** (Activity Selection Problem), then:

```
min_removals = total_intervals − max_non_overlapping
```

Greedy Strategy: Sort intervals by **end time**. Always select the interval that ends earliest among those that don't overlap with the last selected.

Intuition: By picking the interval that ends earliest, we leave the maximum room for future intervals.

Greedy Proof Sketch

Claim: Sorting by end time and greedily selecting non-overlapping intervals is optimal.

Proof by exchange argument:

- Suppose an optimal solution selects interval X with a later end time, when interval Y (earlier end, same start or later) was available.
- Swapping X for Y cannot reduce the number of selected intervals (Y ends earlier, leaving at least as much room).
- Therefore, always picking the earliest-ending non-overlapping interval is at least as good as any other choice.

Implementation

```
def eraseOverlapIntervals(intervals):
    """
    Greedy: sort by end time, count overlapping intervals to remove.
    Time: O(n log n)  |  Space: O(1) (excluding sort)
    """
    if not intervals:
        return 0

    # Sort by end time
    intervals.sort(key=lambda x: x[1])

    max_keep = 1      # We can always keep at least 1 interval
    prev_end = intervals[0][1]   # End time of last kept interval

    for i in range(1, len(intervals)):
        start, end = intervals[i]

        if start >= prev_end:
            # Non-overlapping: keep this interval
            max_keep += 1
            prev_end = end
        # else: overlapping — skip it (implicitly removing it)

    return len(intervals) - max_keep
```

Detailed Trace

Example 1: `[[1,2],[2,3],[3,4],[1,3]]`

Sort by end time:

```
[[1,2], [2,3], [1,3], [3,4]]
```

i	Interval	start	end	prev_end	Overlap?	max_keep
—	[1,2]	—	—	2 (initial)	—	1
1	[2,3]	2	3	2	2 ≥ 2 → keep	2, prev_end = 3
2	[1,3]	1	3	3	1 < 3 → skip (overlaps)	2
3	[3,4]	3	4	3	3 ≥ 3 → keep	3, prev_end = 4

`max_keep = 3, total = 4`
Removals = 4 - 3 = 1 ✓

Example 2: `[[1,2],[1,2],[1,2]]`

Sort by end time:

`[[1,2], [1,2], [1,2]]  (all same)`

i	Interval	start	prev_end	Overlap?	max_keep
—	[1,2]	—	2	—	1
1	[1,2]	1	2	$1 < 2 \rightarrow$ skip	1
2	[1,2]	1	2	$1 < 2 \rightarrow$ skip	1

`max_keep = 1, total = 3`
Removals = 3 - 1 = 2 ✓

Example 3: `[[1,2],[2,3]]`

Sort: `[[1,2],[2,3]]`

i	Interval	start	prev_end	Overlap?	max_keep
—	[1,2]	—	2	—	1
1	[2,3]	2	2	$2 \geq 2 \rightarrow$ keep	2

`max_keep = 2, total = 2`
Removals = 2 - 2 = 0 ✓

Alternative: Sort by Start Time (Wrong!)

```
# Sorting by start time is NOT optimal:
# [[1,10],[2,3],[3,4]] sorted by start = [[1,10],[2,3],[3,4]]
# Greedy picks [1,10] first → removes both [2,3] and [3,4] = 2 removals
# But optimal: remove [1,10], keep [2,3],[3,4] = 1 removal

# Correct answer: sort by END time
```

Test Suite

```
def test_erase_overlap():
    test_cases = [
        ([[1,2],[2,3],[3,4],[1,3]], 1),
        ([[1,2],[1,2],[1,2]], 2),
        ([[1,2],[2,3]], 0),
        ([[-100,-87],[-99,-44],[-98,-19],[-97,-26],[96,99],[95,99],[94,
    ]
```

```
    for intervals, expected in test_cases:
        result = eraseOverlapIntervals([iv[:] for iv in intervals])
        status = "PASS" if result == expected else "FAIL"
        print(f"{status}: erase({intervals[:3]}...) = {result} (expecte

test_erase_overlap()
```

Complexity Analysis

Metric	Value
Time	O(n log n) — dominated by sorting
Space	O(1) — only prev_end and count needed

Key Takeaways

- **Activity Selection Problem:** Classic greedy — sort by end time, greedily pick non-overlapping.
- **Greedy choice:** The interval ending earliest leaves maximum room for subsequent intervals.
- `min_removals = n - max_non_overlapping_intervals`
- Intervals touching at a single point (`start == prev_end`) are **non-overlapping**.
- Sorting by end time (not start time) is essential for the greedy to be correct.

References

- LeetCode #435 — Non-overlapping Intervals
- Cormen et al., *Introduction to Algorithms*, Ch. 16.1 (Activity Selection)
- Skiena, *The Algorithm Design Manual*, Ch. 8.1 (Greedy Algorithms)